



VAO++: Practical Volumetric Ambient Occlusion for Games

Jakub Bokšanský, Adam Pospíšil (Project Wilberforce)
Jiří Bittner (CTU in Prague)

EGSR 19.6.2017

Motivation

- Focus on performance (highly optimized implementation in Unity)
- Unity's built-in SSAO not fast enough and contains artifacts
- VR compatibility (single pass stereo, explicit stereo support)



SSAO

- Adds proximity shadows (enclosed areas, creases, corners...)
- Contributes significantly to perceived realism
- Calculated in screen space
- Limitations:
 - No information about geometry what is not seen
- Advantages
 - Feasible in real time on current GPUs
 - Easy integration with renderers
 - No pre-processing needed
- [Mittring07], [Bavoil08], [Ritschel09]




Scene without effects

VAO++


SSAO Implementation

- For each shaded point, calculate visibility (how much is it obscured by surrounding geometry)
- Surrounding geometry is sampled in normal oriented hemisphere of radius R




Volumetric Ambient Occlusion [Szirmay-Kalos10]

- Transforms hemispherical domain into smaller tangent sphere
 - Uses samples within a disc located in the center of tangent sphere oriented towards camera (2D samples instead of 3D)
 - Each sample represents a line segment inside the sphere in given sampling point
 - Length of the line segment is taken as weight of the sample
 - Samples never overlap



Volumetric Ambient Occlusion II.

- Algorithm calculates where geometry lies wrt. tangent sphere
 - A - Inside (see image)
 - B - Behind (completely unoccluded)
 - C - In front of (completely occluded)
- Visibility (occlusion factor) is calculated as ratio of unoccluded volume to total volume of the tangent sphere




VAO++

- Extensions and improvements to VAO algorithm
- Improvements of visual quality
 - Samples generation
 - Luminance sensitivity
 - Depth range check
- Performance optimizations
 - Adaptive sampling
 - Culling pre-pass


Samples generation

- VAO algorithm requires uniformly distributed samples within a disk
- Poisson distribution is suitable
- Interleaved sampling
 - Size 3x3 - good improvement in quality (number of samples increased virtually 9 times)
 - easy to remove HF noise by low pass filter of matching size
 - Problem with selecting suitable rotation angles
 - Points should not overlap after rotations (prevent sampling of same area)



Samples generation - selecting rotations


- Selecting rotations after generating Poisson distribution yields suboptimal results
 - Goal is to select angles that maximize the distance between positions of rotated samples



(left) rotations selected after generating sample-set. (right) our method


Samples generation - selecting rotations II.

- Take rotations into account when generating Poisson disk distribution
 - We use uniform-step 9 rotations
- Union of sample-set rotated by all angles satisfies Poisson disk distribution requirements




Samples generation - selecting rotations III.

- Final step is to select best permutation of rotated samples
 - Maximize distance between all samples
 - Feasible for 8 samples or less, for larger sample-sets we calculate best permutation for groups of size 8...
- Morton-Ordered for better cache coherence



Samples generation - selecting rotations IV.



Comparison of effect results with sample-set non-optimized (left) vs optimized (right)
Loupe contains enlarged area

Luminance sensitivity I.

- SSAO should be applied on ambient color component before lighting (BRDF) is calculated - not always possible or practical
- VAO++ applies occlusion as final step to rendered image
 - can put occlusion on brightly lit areas, specular highlights, light sources etc.




Scene without effects



VAO++ applied

Luminance sensitivity II.

- To counteract this defect consider luminance of target surface when applying VAO component
- Falloff function to suppress occlusion on bright surfaces
- User can set brightness threshold and slope



Luminance sensitivity III.

- As a luminance use either relative luminance function or Value component of HSV model



Scene without effects




VAO++



VAO++ with luminance sensitivity

Depth Range Check

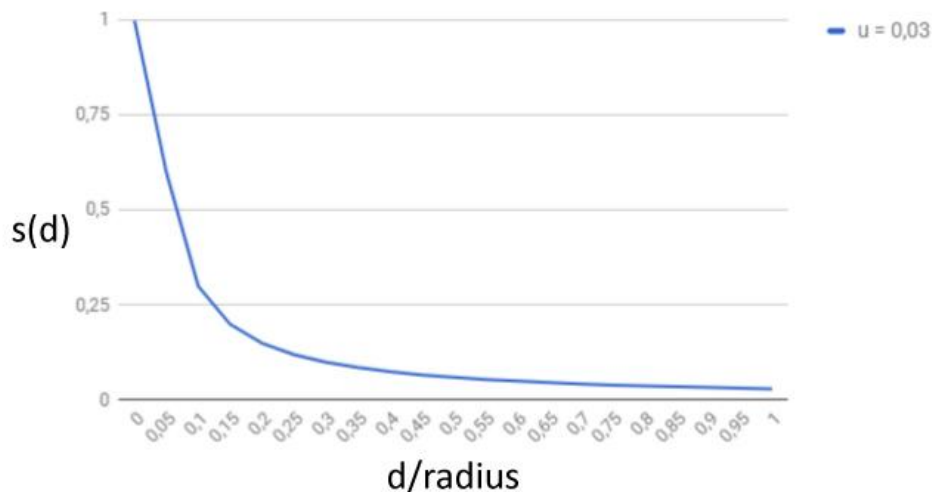
- No geometry information behind what is seen by camera
- Thickness of occluders unknown
 - VAO considers all occluders infinitely thick - leading to unpleasant halo artifact



Depth Range Check II.


- Implement a falloff function describing thickness of geometry seen by camera
 - smoothly decreases thickness with increasing distance
- Resulting occlusion also smoothly decreases around depth discontinuities
- This suppresses halo significantly and works well for dynamic scenes

$$s(d) = \frac{u}{\max(u, \frac{d}{\text{radius}})}$$




Depth Range Check III.

without depth range check



with depth range check




- Extensions and improvements to VAO algorithm
- Improvements of visual quality
 - Samples generation
 - Luminance sensitivity
 - Depth range check
- Performance optimizations
 - Adaptive sampling
 - Culling pre-pass

Adaptive Sampling

- To improve performance use only as many samples as needed
- VAO uses kernel that always displays as a disk on the screen
 - This disk surface shrinks with increasing distance from camera
 - Use fewer samples far away
- Aim is to achieve same density of samples in post-perspective image space


Adaptive Sampling II.

- Pick number of samples to use based on kernel disk surface area after projection onto the screen
- User selects surface area for which lowest number of samples is used
- Algorithm doubles number of samples used when surface area doubles



Adaptive Sampling III.

- Comparison of adaptive sampling on/off



Adaptive On

Adaptive Off

Culling pre-pass

- Areas where no occlusion occurs can be omitted
- Estimate such areas first, calculate VAO only where needed
- Calculate lower quality VAO in a separate pass
- Use result as an estimation

Occlusion estimation texture



Rendered scene



Culling pre-pass II.

- Areas without estimated occlusion could contain undetected details
- Handle culling candidates in one of following ways:
 - Greedy variant - complete omission
 - Careful variant - eight samples regardless of quality settings



Greedy



Careful




Off

Images are
8x zoomed
and gamma corrected

Culling pre-pass III.

- Performance gain scales with resolution
- Especially useful for large screens (4k, VR, etc.)
- Quality loss almost imperceptible (careful variant)




Implementation in Unity


- Image effect attached to camera
 - Inject into pipeline - after opaque rendering pass
 - Read source texture (rendering pipeline output)
 - Write into destination texture
- C# Script + CG/HLSL shader file
- Alternative Command Buffer implementation
- Unity provides depth and normals texture
 - Octa encoded normals + linearized depth packed into 32 bits
- Single Pass Stereo Rendering (VR)
 - Both eyes in single texture
 - Special macros for handling split between eyes
 - Need to calculate effect in world-space

Comparison - 16 SPP, 1920x1080, GeForce GT555M, same radius


VAO++ (48ms) - Speedup 1,25x




Unity's SSAO (60ms)




Comparison



Raytraced AO



VAO++



Unity's SSAO

Comparison

Effect Off



Comparison

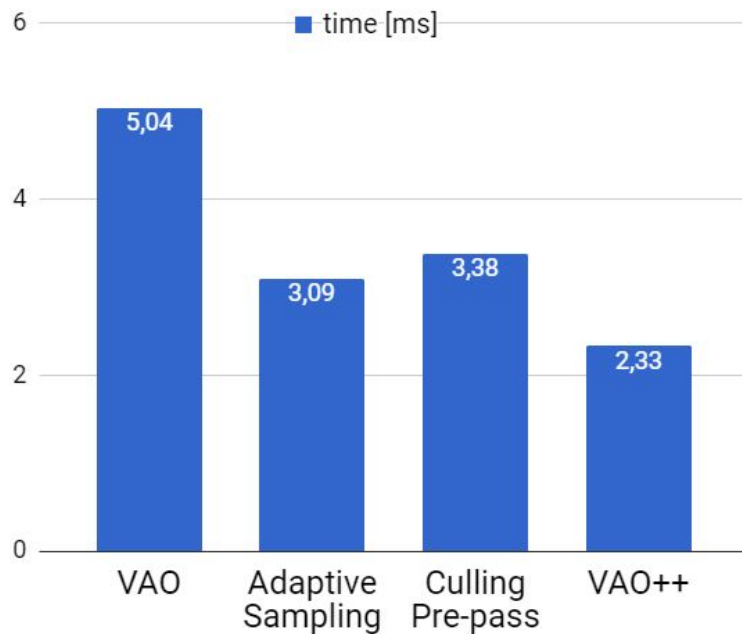
Effect On




Comparison - Video

Performance results

32 SPP 1920x1080



16 SPP 1920x1080



Conclusion

- Samples generation taking rotations for interleaved sampling into account
- Luminance sensitivity and depth range check for more believable results
- Adaptive sampling and culling pre-pass for higher performance with minimal quality loss
- Optimized Unity implementation (available on Asset Store)

Scene by Evgenia



- **Project Wilberforce**
- projectwilberforce.github.io
- projectwilberforce@gmail.com
- Demo available at: <https://projectwilberforce.github.io/vaodemo/>

Thank you!